

Simulation: A Reverse-Engineering Approach to Understand OLS, MLE


Haohan Chen

February 16, 2018


The Big Picture: How Do We Make Predictions?

The setup of a linear model (various notation systems)


```
# -----  
# Simulation for reverse-engineering  
# -----  
  
# Parameters about the simulation  
#-----  
  
# Set N = sample size, k = number of independent variables  
N <- 1000  
k <- 5  
# Set random seed so you get the same result every time.  
set.seed(2.16)  
  
# We first look at the Systematic component  
#-----  
  
# Generate our independent variables X (in a matrix)  
X <- matrix(NA, nrow = N, ncol = k)  
# Note: no requirement about their distribution  
# x1 is a continuous variable drawn from normal dist  
X[, 1] <- rnorm(N, 5, 12)  
plot(density(X[, 1]))
```




```
# x2 is a count variable drawn from a binomial dist, N = 10, p = 0.7  
X[, 2] <- rbinom(N, 10, 0.7)  
hist(X[, 2])
```




```
# x3 is another count variable drawn form a poisson dist, mean = 500
X[, 3] <- rpois(N, 500)
hist(X[, 3])
```



```
# x4 is a proportion variable drawn form a beta distribution
X[, 4] <- rbeta(N, 1, 2)
plot(density(X[, 4]))
```



```
# x5 is a continuous variable drawn form a gamma distribution
X[, 5] <- rgamma(N, 3, 5)
plot(density(X[, 5]))
```



```
# These are our independent variables, but we are not done yet. add the intercept
X <- cbind(rep(1, N), X)
```

```
# What does X look like?
colnames(X) <- c("constant", paste0("x", 1:k))
head(X)
```

```
##      constant      x1 x2  x3      x4      x5
## [1,]         1 -5.762975 6 481 0.02002443 0.2423048
## [2,]         1  7.218190 6 473 0.42308101 1.0173314
## [3,]         1 24.054144 6 500 0.24040910 0.6037380
## [4,]         1 -8.564508 7 497 0.97297753 0.2308831
## [5,]         1  4.036979 8 461 0.29668598 0.1388283
## [6,]         1  6.589043 8 510 0.41919450 0.2561764
```

```
# Now, determine the "ground-truth" parameters
beta <- c(40, 1, 4, 0.2, 9, 20) # Your choice, make sure length of beta = k
names(beta) <- c("(Intercept)", paste0("x", 1:k))
# Then determine the systematic component
sys_component <- X %*% beta
head(sys_component)
```

```
##      [,1]
## [1,] 159.4633
## [2,] 189.9725
## [3,] 202.2926
## [4,] 172.2100
## [5,] 173.6837
## [6,] 189.4853
```

```
# We are done with the systematic component. No assumptions.
# Now we turn to the stochastic component. I am jumping ahead to make the
# Gauss-Markov + normality assumption (not required
# in the first half of demo about OLS, check lecture slides)
```

```
# The Stochastic Component
# -----
epsilon <- rnorm(N, 0, 2) # mean = 0, sd = 3
stochastic_component <- epsilon
```

Getting OLS estimators for linear models

(Whiteboard demo)

```
# -----  
# Put together a simulated dataset for linear models  
# -----  
  
# Y for linear relation  
# -----  
Y <- sys_component + stochastic_component  
  
# Voila: we are done simulating for a linear model. Think about this  
# We observe only Y and X, while not observing beta, epsilon. like this  
data_linear <- data.frame(Y = Y, X)  
head(data_linear)
```

##	Y	constant	x1	x2	x3	x4	x5
## 1	156.2328	1	-5.762975	6	481	0.02002443	0.2423048
## 2	188.9919	1	7.218190	6	473	0.42308101	1.0173314
## 3	201.6978	1	24.054144	6	500	0.24040910	0.6037380
## 4	172.2819	1	-8.564508	7	497	0.97297753	0.2308831
## 5	173.5673	1	4.036979	8	461	0.29668598	0.1388283
## 6	189.1417	1	6.589043	8	510	0.41919450	0.2561764

```
# I want to separate out a few samples to demo prediction so I just  
# Get a slice out of the complete dataset  
# You heard about "training" and "test" set. But let's not discuss it now  
data_new <- data_linear[1:10, ]  
data_fit <- data_linear[-(1:10), ]  
head(data_new)
```

##	Y	constant	x1	x2	x3	x4	x5
## 1	156.2328	1	-5.762975	6	481	0.02002443	0.2423048
## 2	188.9919	1	7.218190	6	473	0.42308101	1.0173314
## 3	201.6978	1	24.054144	6	500	0.24040910	0.6037380
## 4	172.2819	1	-8.564508	7	497	0.97297753	0.2308831
## 5	173.5673	1	4.036979	8	461	0.29668598	0.1388283
## 6	189.1417	1	6.589043	8	510	0.41919450	0.2561764

```
head(data_fit) # we use this to fit the model
```

##	Y	constant	x1	x2	x3	x4	x5
## 11	191.2679	1	10.0118090	6	519	0.455841528	0.5060738
## 12	197.2827	1	16.7810333	9	487	0.047185283	0.4537083
## 13	176.7305	1	0.2876557	7	486	0.108345348	0.5818760
## 14	164.7003	1	-7.4760277	4	532	0.007592964	0.5590694
## 15	204.2397	1	26.3867475	8	497	0.412413849	0.1984410
## 16	157.4544	1	-22.7328290	7	502	0.308822018	0.5325092

```
# -----  
# Get OLS estimators for the simulated data  
# -----  
  
# Fit the model  
# -----  
# You've done it a million times with R function. I suppose
```

```

m_ols <- lm(Y ~ x1 + x2 + x3 + x4 + x5, data = data_fit)
coef(m_ols) # The estimated coefficient

## (Intercept)          x1          x2          x3          x4          x5
## 41.9894037    1.0094391    4.0234648    0.1953583    8.8127729    20.0088841

beta # The "ground truth"

## (Intercept)          x1          x2          x3          x4          x5
##      40.0           1.0           4.0           0.2           9.0          20.0

# You find from above the estimated beta is almost the "ground truth"
# So you know it's working

# Inference (uncertainty of beta's)
-----
# Variance in estimation?
vcov(m_ols)

##              (Intercept)              x1              x2              x3
## (Intercept) -2.1840170426 -3.216162e-04  1.304090e-02  4.144213e-03
## x1          -0.0003216162 -2.857943e-05 -9.748134e-06  1.124962e-06
## x2           0.0130408975 -9.748134e-06 -2.077069e-03  2.268167e-06
## x3           0.0041442129  1.124962e-06  2.268167e-06 -8.329995e-06
## x4           0.0225956737 -1.208696e-05  6.535147e-04 -6.172582e-06
## x5           0.0106067288 -1.070876e-05  2.058067e-04  1.631596e-05
##              x4              x5
## (Intercept)  2.259567e-02  1.060673e-02
## x1          -1.208696e-05 -1.070876e-05
## x2           6.535147e-04  2.058067e-04
## x3          -6.172582e-06  1.631596e-05
## x4          -7.455160e-02  1.484871e-03
## x5           1.484871e-03 -3.370285e-02

# Simulate beta to get confidence interval of your coefficients
N_sim <- 10000 # note: a different thing from N defined above
library(MASS) # library for the mvrnorm function
beta_sim <- mvrnorm(N_sim, mu = coef(m_ols), Sigma = vcov(m_ols))
head(beta_sim)

##      (Intercept)          x1          x2          x3          x4          x5
## [1,]  41.45633    1.002515    4.043962    0.1962681    9.275595    19.67583
## [2,]  45.40551    1.011187    3.997409    0.1887883    8.841115    19.95738
## [3,]  43.83916    1.011979    4.060883    0.1909360    8.621648    20.13476
## [4,]  41.46801    1.016323    4.059884    0.1965087    8.832222    19.46800
## [5,]  42.59391    1.011847    4.031356    0.1938568    8.795301    20.06492
## [6,]  40.00121    1.015074    4.047794    0.1987134    8.784852    20.18626

# Get summary statistics of this sample
beta_confint95 <- apply(beta_sim, 2, function(x) quantile(x, c(0.025, 0.5, 0.975)))
beta_confint95

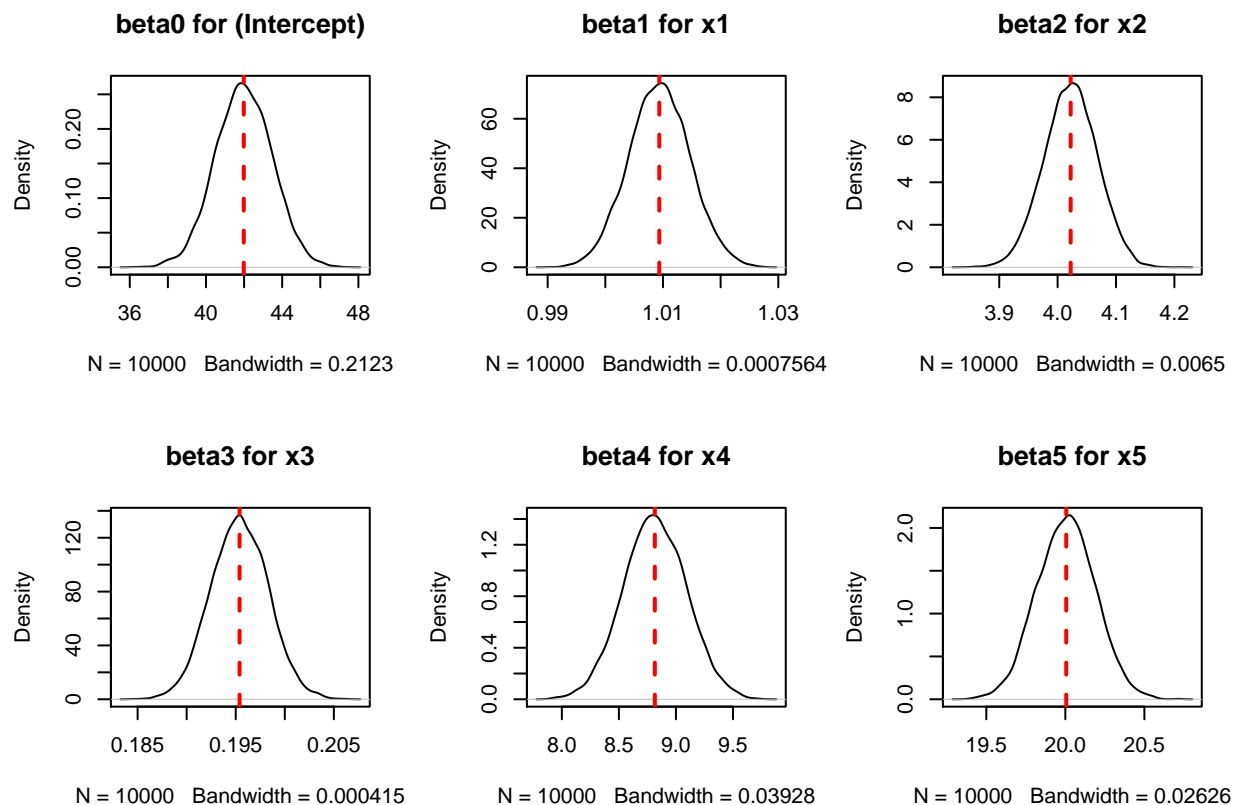
##      (Intercept)          x1          x2          x3          x4          x5
## 2.5%    39.08532    0.9990038    3.932624    0.1897456    8.278858    19.64569
## 50%     41.97990    1.0093562    4.022835    0.1953956    8.813528    20.00757
## 97.5%   44.89896    1.0199328    4.110426    0.2011247    9.362238    20.36326

```


```
# Plot the distribution of beta's
par(mfrow = c(2, ceiling((k+1)/2)))
for (i in 1:ncol(beta_sim)){
  plot(density(beta_sim[, i]),
       main = paste(paste0("beta", i-1), "for", colnames(beta_sim)[i]))
  abline(v = mean(beta_sim[, i]), col = "red", lty = 2, lwd = 2)
}
# Plot coefficients (my way... feel free to do it in another way)
beta_confint95_2 <- as.data.frame(t(beta_confint95))
beta_confint95_2$variable <- row.names(beta_confint95_2)
names(beta_confint95_2)
```

```
## [1] "2.5%"      "50%"       "97.5%"     "variable"
```

```
library(ggplot2)
```



```
ggplot(beta_confint95_2, aes(y = `50%`, x = variable)) + geom_point() +
  geom_linerange(aes(ymin = `2.5%`, ymax = `97.5%`))
```



```
# Some are very small, because not at the same scale.

# Prediction (uncertainty about a new y)
# -----
# Task: you have a new x (vector), you want to know what the corresponding y is
head(data_new)

##          Y constant      x1 x2  x3          x4          x5
## 1 156.2328         1 -5.762975  6 481 0.02002443 0.2423048
## 2 188.9919         1  7.218190  6 473 0.42308101 1.0173314
## 3 201.6978         1 24.054144  6 500 0.24040910 0.6037380
## 4 172.2819         1 -8.564508  7 497 0.97297753 0.2308831
## 5 173.5673         1  4.036979  8 461 0.29668598 0.1388283
## 6 189.1417         1  6.589043  8 510 0.41919450 0.2561764

# Let's take one of them
x_new <- data_new[1, c("constant", paste0("x", 1:k))]
x_new

##   constant      x1 x2  x3          x4          x5
## 1         1 -5.762975  6 481 0.02002443 0.2423048

# Challenge: What one of the x to be certain value?

# Want: y_new ~ N(mu_y_new, sigma_y).
# When we predict, we get a distribution, not one number

# First we get mu_y_new: harder part
```

```

mu_y_new <- as.matrix(x_new, byrow = T) %*% t(beta_sim)
dim(mu_y_new) # It's a vector of 10000 sample of predicted y

## [1]      1 10000

mu_y_new <- as.vector(mu_y_new) # just to clean up. transform it into a vector
mu_y_new[1:30] # See how it looks like

## [1] 159.3009 159.3825 159.2640 159.3850 159.2339 159.0864 159.6307
## [8] 159.5388 159.2754 159.6779 159.3888 159.0747 159.2673 159.2374
## [15] 159.4132 159.0057 159.1771 159.1440 159.3516 159.4732 159.1585
## [22] 159.3798 159.2244 159.5618 159.1205 159.1474 159.0909 159.5011
## [29] 159.2602 158.9382

# Then we get sigma_y. easy
sigma_y <- sigma(m_ols)
# Want it by hand (using week 2 page 11 formula)
sigma_y <- sqrt((t(residuals(m_ols)) %*% residuals(m_ols)) /
  (N - (k + 1)))

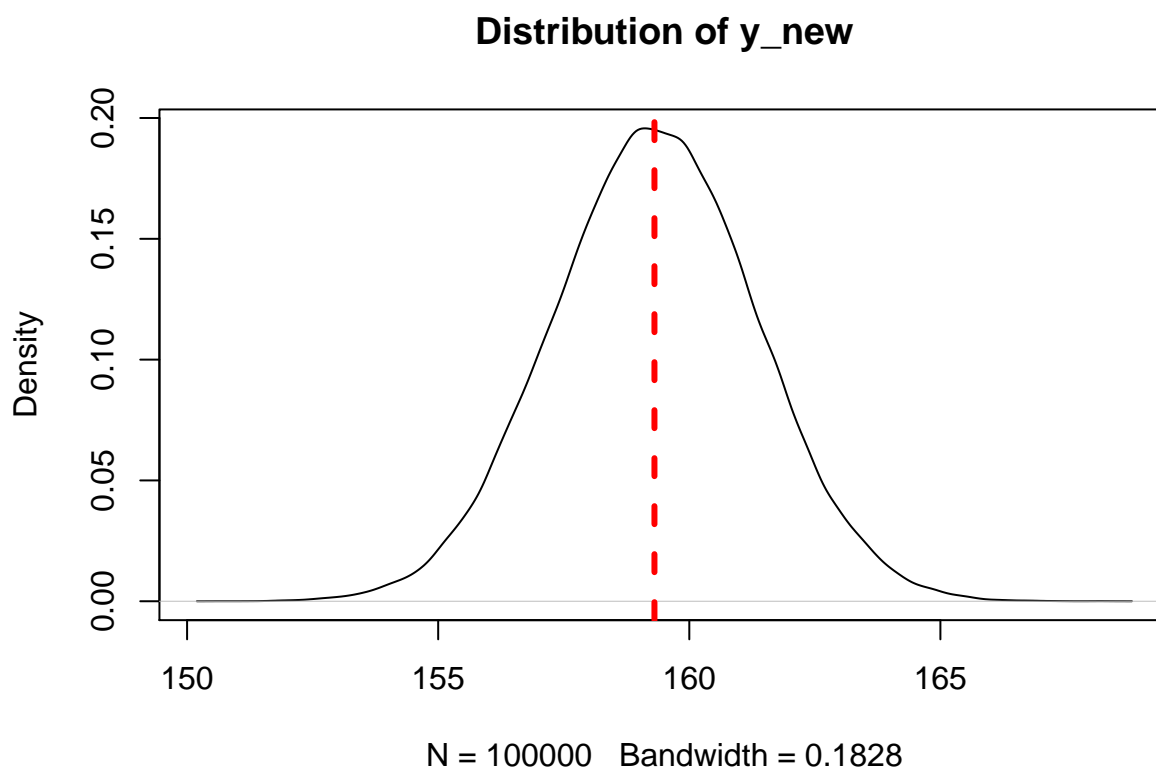
# Then, simulate y_new useing this list of mu_y_new
y_new <- rnorm(100000, mu_y_new, sigma_y)

# We are done with prediction! You have predicted y_new. Plot it, summarise it
quantile(y_new, c(0.025, 0.5, 0.975)) # 95% confidence interval

##      2.5%      50%      97.5%
## 155.3163 159.3115 163.2921

par(mfrow = c(1, 1))
plot(density(y_new), main = "Distribution of y_new")
abline(v = mean(y_new), col = "red", lty = 2, lwd = 3)

```

```
# Expected new y? Simply take the mean!
mean(y_new) # this is your expected value
```

```
## [1] 159.3062
```

```
# Now, challenge: do it without using lm().
est_beta_manual <- solve(t(X) %*% X) %*% t(X) %*% Y
est_beta_manual
```

```
##           [,1]
## constant 41.8937720
## x1       1.0093564
## x2       4.0231604
## x3       0.1955287
## x4       8.8384388
## x5      20.0106527
```

```
e <- Y - X %*% est_beta_manual
# a bit different because we remove 10 cases for the prediction task
est_beta_vcov <- as.numeric((t(e) %*% e) / (N - (k + 1))) * solve(t(X) %*% X)
est_beta_vcov
```

```
##           constant           x1           x2           x3
## constant 2.1561203113 3.240649e-04 -1.286443e-02 -4.092776e-03
## x1       0.0003240649 2.815450e-05 9.242122e-06 -1.122038e-06
## x2      -0.0128644328 9.242122e-06 2.052472e-03 -2.295528e-06
## x3      -0.0040927759 -1.122038e-06 -2.295528e-06 8.229755e-06
## x4      -0.0217919051 1.907055e-05 -6.444249e-04 5.063731e-06
```

```
## x5          -0.0102673796  9.309850e-06 -1.908965e-04 -1.644516e-05
##              x4              x5
## constant -2.179191e-02 -1.026738e-02
## x1         1.907055e-05  9.309850e-06
## x2        -6.444249e-04 -1.908965e-04
## x3         5.063731e-06 -1.644516e-05
## x4         7.320729e-02 -1.365543e-03
## x5        -1.365543e-03  3.316061e-02
```

MLE estimator for linear model

(Whiteboard demo)

```
# Fit the model
m_mle <- glm(Y ~ x1 + x2 + x3 + x4 + x5, data = data_fit, family = "gaussian")
# Get estimated beta
coef(m_mle)
```

```
## (Intercept)          x1          x2          x3          x4          x5
## 41.9894037    1.0094391    4.0234648    0.1953583    8.8127729    20.0088841
```

beta

```
## (Intercept)          x1          x2          x3          x4          x5
##         40.0          1.0          4.0          0.2          9.0         20.0
```

```
# Get variance of the estimation
vcov(m_mle)
```

```
##              (Intercept)          x1          x2          x3
## (Intercept)  2.1840170426  3.216162e-04 -1.304090e-02 -4.144213e-03
## x1           0.0003216162  2.857943e-05  9.748134e-06 -1.124962e-06
## x2          -0.0130408975  9.748134e-06  2.077069e-03 -2.268167e-06
## x3          -0.0041442129 -1.124962e-06 -2.268167e-06  8.329995e-06
## x4          -0.0225956737  1.208696e-05 -6.535147e-04  6.172582e-06
## x5          -0.0106067288  1.070876e-05 -2.058067e-04 -1.631596e-05
##              x4              x5
## (Intercept) -2.259567e-02 -1.060673e-02
## x1           1.208696e-05  1.070876e-05
## x2          -6.535147e-04 -2.058067e-04
## x3           6.172582e-06 -1.631596e-05
## x4           7.455160e-02 -1.484871e-03
## x5          -1.484871e-03  3.370285e-02
```

```
# Inference
# -----
# (exercise)
```

```
# Predictions
# -----
# (exercise)
```

Modeling Dicotomous Outcome Using Generalized Linear Model: Setup

(Whiteboard demonstration)

```
rm(list=ls())


# Just copy pasting. should've made it a function. pressed for time.

# Parameters about the simulation (nothing to do with the model)
#-----


# Set N = sample size, k = number of independent variables
N <- 10000
k <- 5
# Set random seed so you get the same result every time.
set.seed(2.16)

# We first look at the Systematic component
#-----


# Generate our independent variables X (in a matrix)
X <- matrix(NA, nrow = N, ncol = k)
# Note: no requirement about their distribution
# x1 is a continuous variable drawn from normal dist
X[, 1] <- rnorm(N, 2, 12)
plot(density(X[, 1]))
```



```
# x2 is a count variable drawn from a binomial dist, N = 10, p = 0.7
X[, 2] <- rbinom(N, 3, 0.7)
hist(X[, 2])
```




```
# x3 is another count variable drawn form a poisson dist, mean = 500  
X[, 3] <- rpois(N, 2)  
hist(X[, 3])
```



```
# x4 is a proportion variable drawn form a beta distribution  
X[, 4] <- rbeta(N, 1, 2)  
plot(density(X[, 3]))
```


density.default(x = X[, 3])



N = 10000 Bandwidth = 0.2047

```
# x5 is a continuous variable drawn form a gamma distribution  
X[, 5] <- rgamma(N, 3, 5)  
plot(density(X[, 4]))
```

density.default(x = X[, 4])



N = 10000 Bandwidth = 0.03382

```
# These are our independent variables, but we are not done yet. add the intercept
X <- cbind(rep(1, N), X)
```

```
# What does X look like?
colnames(X) <- c("constant", paste0("x", 1:k))
head(X)
```

```
##      constant      x1 x2 x3      x4      x5
## [1,]         1 -8.762975  1  1 0.14158054 0.2043488
## [2,]         1  4.218190  1  2 0.35102508 0.8078260
## [3,]         1 21.054144  3  2 0.02282631 0.6614940
## [4,]         1 -11.564508  2  3 0.25226094 1.5345986
## [5,]         1  1.036979  1  2 0.18886517 0.4645816
## [6,]         1  3.589043  1  0 0.35007593 1.1798757
```

```
# Now, determine the "ground-truth" parameters
beta <- c(0.02, 0.01, 0.04, 0.02, 0.09, 0.05) # Your choice, make sure length of beta = k
names(beta) <- c("(Intercept)", paste0("x", 1:k))
# Then determine the systematic component
sys_component <- X %%% beta
head(sys_component)
```

```
##           [,1]
## [1,] 0.01532995
## [2,] 0.21416546
## [3,] 0.42567051
## [4,] 0.14378834
```

```
## [5,] 0.15059674
## [6,] 0.18639105

# We are done with the systematic component. No assumptions.
# Now we turn to the stochastic component. I am jumping a head to make the
# Gauss-Markov + normality assumption (not required
# in the first half of demo about OLS, check lecture slides)

# The Stochastic Component
# -----
epsilon <- rnorm(N, 0, 0.001) # mean = 0, sd = 3
stochastic_component <- epsilon

# -----
# Set up the simulated dataset
# -----

# The probabilities with a logistic link
pi_logit <- 1 / (1 + exp(-sys_component))
# The probabilities with a probit link
pi_probit <- pnorm(sys_component)

set.seed(1)
# The outcome y of a logit link
y_logit <- sapply(pi_logit, function(x) rbinom(1, 1, x))
table(y_logit)

## y_logit
##      0      1
## 4377 5623

# The outcome y of a probit link
y_probit <- sapply(pi_probit, function(x) rbinom(1, 1, x))
table(y_probit)

## y_probit
##      0      1
## 4149 5851

# Dataset
data_logit <- data.frame(Y = y_logit, X)
data_probit <- data.frame(Y = y_probit, X)
```

MLE Estimator (Logit)

```
# -----
# Fit the model
# -----
m_logit <- glm(Y ~ x1 + x2 + x3 + x4 + x5, data = data_logit, family = binomial(link = "logit"))
coef(m_logit)

## (Intercept)          x1          x2          x3          x4          x5
## -0.02225989  0.01058423  0.07264404  0.01315730  0.10814500  0.05971384
beta
```



```
## (Intercept)      x1      x2      x3      x4      x5
##      0.02      0.01      0.04      0.02      0.09      0.05

# -----
# Inference (exercise)
# -----

# -----
# Prediction (exercise)
# -----
```

MLE Estimator (Probit)

```
# -----
# Fit the model
# -----
m_probit <- glm(Y ~ x1 + x2 + x3 + x4 + x5, data = data_logit, family = binomial(link = "probit"))
coef(m_probit)

## (Intercept)      x1      x2      x3      x4
## -0.013643820  0.006600958  0.045366949  0.008194967  0.067244283
##      x5
##  0.037436864

beta

## (Intercept)      x1      x2      x3      x4      x5
##      0.02      0.01      0.04      0.02      0.09      0.05

vcov(m_probit)

##      (Intercept)      x1      x2      x3
## (Intercept)  2.358841e-03 -1.287707e-06 -5.231372e-04 -1.498670e-04
## x1          -1.287707e-06  1.110506e-06 -1.275007e-07  4.897661e-09
## x2          -5.231372e-04 -1.275007e-07  2.514135e-04 -1.534684e-06
## x3          -1.498670e-04  4.897661e-09 -1.534684e-06  7.730016e-05
## x4          -9.311772e-04 -5.034428e-07 -8.383379e-06 -9.198598e-06
## x5          -8.067830e-04 -8.863650e-07  1.577707e-06  5.804393e-07
##      x4      x5
## (Intercept) -9.311772e-04 -8.067830e-04
## x1          -5.034428e-07 -8.863650e-07
## x2          -8.383379e-06  1.577707e-06
## x3          -9.198598e-06  5.804393e-07
## x4          2.834946e-03  3.863221e-05
## x5          3.863221e-05  1.315015e-03

# -----
# Inference (exercise)
# -----

# -----
# Prediction (exercise)
# -----
```